

Time Series HW 3

Kenny Flagg and Paul Harmon

who have not yet worked together and want the bonus!

September 16, 2016

HW 3

You can alone or in pairs that may or may not be people you worked with before. You can discuss it with old partners but should try work as much as possible with new collaborators. 5% bonus if you find someone completely new to work with - that you did not work with on first two assignments.

I mentioned de-seasonalizing of time series, where the seasonal variation is removed from the series to highlight variation at either higher or lower frequencies. There are a variety of techniques for doing this but the simplest is to just subtract the mean for each month from the observations. And the easiest way to find that is using `lm(y~month, data=...)`.


Note that Kenny and Paul have *NOT* worked together on any previous homework assignments.

- 1) *For the Bozeman temperature data from HW 1 and 2, estimate a model with month only, subtract its fitted values from the responses (or just extract residuals). Plot the residuals vs the fractional Year variable and compare the plot of this result to the plot of the original time series.*

```
weather = read.csv("rawbozemadata.csv", header = TRUE)

weather$YEAR = substr(as.character(weather$DATE), 0, nchar(weather$DATE)-2)
weather$MONTH = as.factor(substr(as.character(weather$DATE), 5, nchar(weather$DATE)))
weather$year = (floor(weather$DATE/100))
weather$Month<-round((weather$DATE/100-weather$year)*100,1)
weather$Yearfrac<-weather$year+(weather$Month-1)/12 #makes it an integer variable
#fit the linear model
lm.weather = lm(MMXT ~ MONTH, data = weather)
#lm.weather$residuals
plot(weather$Yearfrac, lm.weather$residuals, pch = 20, col = "orange4",
      main = "Residuals vs. Year", xlab = "Year", ylab = "Residuals")
```

Residuals vs. Year



- 2) In the de-seasonalized Bozeman temperature data set, re-assess evidence for the linear trend. Compare the result (test statistic, degrees of freedom and size of p-value) of just fitting a linear time trend in these de-seasonalized responses to the result from our original model with a linear year component and a month adjustment (not the quadratic trend model).

```
lm.seasoned <- lm(MMXT ~ Yearfrac + MONTH, data = weather)
lm.deseasoned <- lm(lm.weather$residuals ~ Yearfrac, data = weather)
summary(lm.seasoned)
```

```
##
## Call:
## lm(formula = MMXT ~ Yearfrac + MONTH, data = weather)
##
## Residuals:
```

| | Min | 1Q | Median | 3Q | Max |
|--|----------|---------|--------|--------|---------|
| | -20.5022 | -2.9005 | 0.1112 | 3.0412 | 12.6950 |

```
##
## Coefficients:
```

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|------------|------------|---------|--------------|
| (Intercept) | -69.290162 | 7.266158 | -9.536 | < 2e-16 *** |
| Yearfrac | 0.051674 | 0.003706 | 13.942 | < 2e-16 *** |
| MONTH02 | 3.701446 | 0.604984 | 6.118 | 1.23e-09 *** |
| MONTH03 | 11.189381 | 0.604985 | 18.495 | < 2e-16 *** |
| MONTH04 | 21.977317 | 0.604986 | 36.327 | < 2e-16 *** |
| MONTH05 | 31.224003 | 0.606297 | 51.500 | < 2e-16 *** |
| MONTH06 | 39.708393 | 0.606299 | 65.493 | < 2e-16 *** |
| MONTH07 | 49.564963 | 0.608994 | 81.388 | < 2e-16 *** |
| MONTH08 | 48.465835 | 0.607645 | 79.760 | < 2e-16 *** |
| MONTH09 | 37.629365 | 0.608976 | 61.791 | < 2e-16 *** |
| MONTH10 | 25.822126 | 0.607627 | 42.497 | < 2e-16 *** |

```
## MONTH11      10.315109    0.607657   16.975 < 2e-16 ***
## MONTH12       1.732846    0.608989    2.845  0.0045 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.597 on 1361 degrees of freedom
## Multiple R-squared:  0.9349, Adjusted R-squared:  0.9343
## F-statistic: 1628 on 12 and 1361 DF,  p-value: < 2.2e-16
```

```
summary(lm.deseasoned)
```

```
##
## Call:
## lm(formula = lm.weather$residuals ~ Yearfrac, data = weather)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -20.4490  -2.9043   0.0912   3.0355  12.7135
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.012e+02  7.229e+00  -13.99  <2e-16 ***
## Yearfrac      5.166e-02  3.691e-03   14.00  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.579 on 1372 degrees of freedom
## Multiple R-squared:  0.1249, Adjusted R-squared:  0.1243
## F-statistic: 195.9 on 1 and 1372 DF,  p-value: < 2.2e-16
```

After de-seasonalization, there is strong evidence of a linear trend in MMXT through time ($t_{1372} = 14.00$, $p\text{-value} < 0.0001$). This is very similar to the result when when month is included as a factor, where the estimated year term has $t_{1361} = 13.94$ with $p\text{-value} < 0.0001$. It seems like the degrees of freedom for the de-seasonalized model is wrong because it doesn't account for estimating the monthly means.

- 3) *I briefly discussed the HADCRUT data set in class. We will consider the HADCRUT4 series of temperature anomalies for the Northern Hemisphere. The fully up to date data set is available at: http://www.metoffice.gov.uk/hadobs/hadcrut4/data/current/time_series/HadCRUT.4.4.0.0.monthly_nh.txt*

Download the ensemble median monthly northern hemisphere temperature data. We will use the entire time series that is currently available (January 1850 to July 2016). You might want to read http://www.metoffice.gov.uk/hadobs/hadcrut4/data/current/series_format.html for more information on the columns in the data set.

Because the time series is complete over the time frame under consideration, you can use `ts()` to create a time variable instead of messing around with their `Year/Month` variable.

Make a plot versus time of the ensemble medians and use that as your response variable in the following questions. Discuss trend, seasonality, outliers, and variability.

```
hadcrut <- read.table('HadCRUT.4.4.0.0.monthly_nh.txt',
                     col.names = c('Date', 'median.temp.anom',
                                   'bias.l95', 'bias.u95',
                                   'meas.l95', 'meas.u95',
                                   'coverage.l95', 'coverage.u95',
                                   'error.l95', 'error.u95',
```

```

                                'l95', 'u95'))
hadcrut.ts <- ts(hadcrut$median.temp.anom, start = 1850, frequency = 12)

```

- 4) Our main focus with these data will be on estimating the long-term trend, starting with polynomial trend models. But first, check for seasonality in a model that accounts for a linear time trend. Use our previous fractional year for the trend. Report an *effects* plot and a test for the month model component.

Note: when you use `time()` to generate the `Year` variable from a time series object it retains some time series object information that can cause conflicts later on. Create a new variable in your data.frame that uses something like `as.vector(time(tsdataname))`.

```

hadcrut$yearfrac <- as.vector(time(hadcrut.ts))
hadcrut$month <- factor(round(1 + 12 * (hadcrut$yearfrac %% 1), 0))
had.lm1 <- lm(median.temp.anom ~ month + yearfrac, data = hadcrut)
summary(had.lm1)

```

```

##
## Call:
## lm(formula = median.temp.anom ~ month + yearfrac, data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.24294 -0.19385 -0.02191  0.18118  1.15426
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.027e+01  2.590e-01 -39.654  < 2e-16 ***
## month2      -3.281e-02  3.141e-02  -1.045  0.296288
## month3      -1.072e-01  3.141e-02  -3.413  0.000656 ***
## month4      -1.621e-02  3.141e-02  -0.516  0.605890
## month5      -5.031e-03  3.141e-02  -0.160  0.872754
## month6       5.062e-02  3.141e-02   1.612  0.107154
## month7       7.306e-02  3.141e-02   2.326  0.020102 *
## month8       7.620e-02  3.145e-02   2.423  0.015502 *
## month9       4.339e-02  3.145e-02   1.379  0.167953
## month10      3.006e-02  3.145e-02   0.956  0.339359
## month11     -6.260e-02  3.145e-02  -1.990  0.046701 *
## month12     -5.818e-02  3.145e-02  -1.850  0.064507 .
## yearfrac     5.285e-03  1.335e-04  39.588  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.287 on 1986 degrees of freedom
## Multiple R-squared:  0.4525, Adjusted R-squared:  0.4492
## F-statistic: 136.8 on 12 and 1986 DF, p-value: < 2.2e-16

```

```

library(car)
Anova(had.lm1, test = "F", type = "III") #test for the month component

```

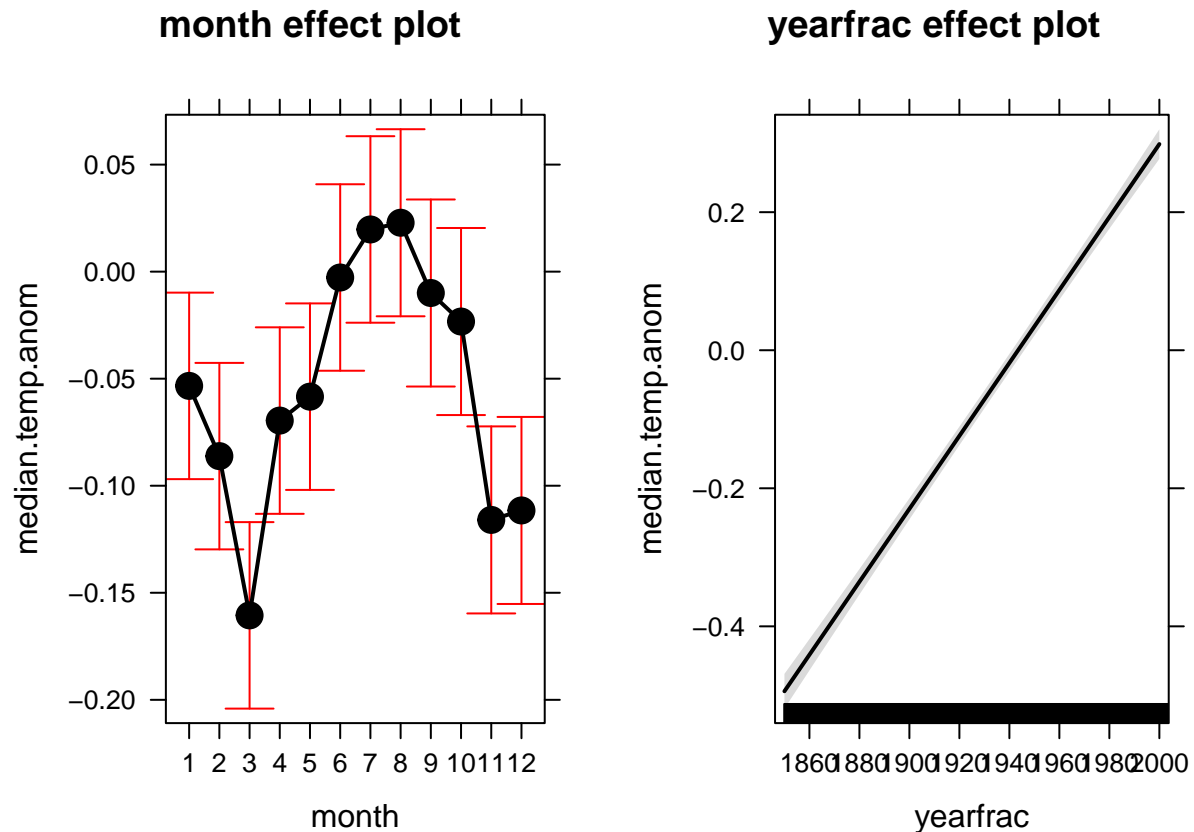
```

## Anova Table (Type III tests)
##
## Response: median.temp.anom

```

```
##           Sum Sq   Df F value    Pr(>F)
## (Intercept) 129.517     1 1572.417 < 2.2e-16 ***
## month        6.104    11   6.737 3.389e-11 ***
## yearfrac     129.090     1 1567.241 < 2.2e-16 ***
## Residuals    163.583 1986
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
require(effects)
plot(allEffects(had.lm1))
```




- 5) Check the residuals versus fitted values for any evidence of nonlinearity in the residuals vs fitted that was missed by the model with a linear trend and month component. Also note any potential issues with the constant variance assumption.

Based on the plot below, there appears to be a wavy pattern occurring in the residuals over time that is not being accounted for in the model with a linear trend and month component. Furthermore, the spread of observations is not constant; it appears that the spread for smaller fitted values is wider near fitted values of -0.5 and much narrower for larger fitted values.

```
plot(had.lm1$fitted.values, had.lm1$residuals, pch = 20, col = "blue2", main = "Residuals vs. Fitted")
```

Residuals vs. Fitted Values



- 6) You can add higher order polynomial terms to models using $x1 + I(x1^2) + I(x1^3) \dots$ or using the `poly` function, such as `poly(x1,3,raw=T)` for a cubic polynomial that includes the linear and quadratic components (we want this!). The `raw=T` keeps the variables in their raw or input format. Estimate the same model but now using polynomial trends that steps up from linear (`poly(time,1,raw=T)`) and stop when you get a failure to estimate a part of the model. Briefly discuss what happened.

The following code estimates the i th-degree polynomial, where i ranges from 1 to 10. Thus, we see that from the returned summary output we get estimates for the coefficients up to the 5th degree polynomial. (Note that had we successfully estimated every part of the model for all 10 coefficients, we would have simply expanded our for-loop to contain more iterations.) It appears that the 5th degree

```
for (i in 1:10){
  had.lm.poly = lm(median.temp.anom ~ poly(yearfrac,i,raw = TRUE), data = hadcrut)
  print(summary(had.lm.poly))
}
```

```
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.2422 -0.1955 -0.0159  0.1787  1.1222
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -1.027e+01  2.622e-01  -39.17  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)  5.285e-03  1.356e-04   38.98  <2e-16 ***
```

```

## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2915 on 1997 degrees of freedom
## Multiple R-squared:  0.4321, Adjusted R-squared:  0.4318
## F-statistic: 1519 on 1 and 1997 DF,  p-value: < 2.2e-16
##
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.24901 -0.16217  0.01274  0.16489  1.00989
##
## Coefficients:
##                                Estimate Std. Error t value Pr(>|t|)
## (Intercept)                   2.222e+02  1.057e+01   21.02  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)1 -2.353e-01  1.094e-02  -21.52  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)2  6.223e-05  2.829e-06   22.00  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2616 on 1996 degrees of freedom
## Multiple R-squared:  0.5429, Adjusted R-squared:  0.5425
## F-statistic: 1186 on 2 and 1996 DF,  p-value: < 2.2e-16
##
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.24329 -0.16121  0.00362  0.16110  1.00864
##
## Coefficients:
##                                Estimate Std. Error t value Pr(>|t|)
## (Intercept)                  -2.895e+03  4.785e+02  -6.051 1.72e-09 ***
## poly(yearfrac, i, raw = TRUE)1  4.606e+00  7.430e-01   6.199 6.88e-10 ***
## poly(yearfrac, i, raw = TRUE)2 -2.443e-03  3.844e-04  -6.355 2.58e-10 ***
## poly(yearfrac, i, raw = TRUE)3  4.319e-07  6.628e-08   6.517 9.08e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2589 on 1995 degrees of freedom
## Multiple R-squared:  0.5525, Adjusted R-squared:  0.5518
## F-statistic: 820.9 on 3 and 1995 DF,  p-value: < 2.2e-16
##
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)

```

```

##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.22121 -0.14956  0.00567  0.15313  1.06170
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.144e+05  2.149e+04   9.978  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)1 -4.454e+02  4.449e+01 -10.010  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)2  3.469e-01  3.454e-02  10.043  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)3 -1.201e-04  1.191e-05 -10.078  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)4  1.558e-08  1.540e-09  10.115  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2526 on 1994 degrees of freedom
## Multiple R-squared:  0.5743, Adjusted R-squared:  0.5735
## F-statistic: 672.5 on 4 and 1994 DF,  p-value: < 2.2e-16
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.22121 -0.14956  0.00567  0.15313  1.06170
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.144e+05  2.149e+04   9.978  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)1 -4.454e+02  4.449e+01 -10.010  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)2  3.469e-01  3.454e-02  10.043  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)3 -1.201e-04  1.191e-05 -10.078  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)4  1.558e-08  1.540e-09  10.115  <2e-16 ***
## poly(yearfrac, i, raw = TRUE)5           NA           NA         NA         NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2526 on 1994 degrees of freedom
## Multiple R-squared:  0.5743, Adjusted R-squared:  0.5735
## F-statistic: 672.5 on 4 and 1994 DF,  p-value: < 2.2e-16
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.22121 -0.14956  0.00567  0.15313  1.06170
##
## Coefficients: (2 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)

```



```

## (Intercept)                2.144e+05  2.149e+04   9.978   <2e-16 ***
## poly(yearfrac, i, raw = TRUE)1 -4.454e+02  4.449e+01 -10.010   <2e-16 ***
## poly(yearfrac, i, raw = TRUE)2  3.469e-01  3.454e-02  10.043   <2e-16 ***
## poly(yearfrac, i, raw = TRUE)3 -1.201e-04  1.191e-05 -10.078   <2e-16 ***
## poly(yearfrac, i, raw = TRUE)4  1.558e-08  1.540e-09  10.115   <2e-16 ***
## poly(yearfrac, i, raw = TRUE)5          NA          NA          NA          NA
## poly(yearfrac, i, raw = TRUE)6          NA          NA          NA          NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2526 on 1994 degrees of freedom
## Multiple R-squared:  0.5743, Adjusted R-squared:  0.5735
## F-statistic: 672.5 on 4 and 1994 DF,  p-value: < 2.2e-16
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.24480 -0.15272  0.00679  0.15599  1.01582
##
## Coefficients: (2 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -3.950e+06  7.029e+05  -5.620 2.18e-08 ***
## poly(yearfrac, i, raw = TRUE)1  9.617e+03  1.698e+03   5.663 1.70e-08 ***
## poly(yearfrac, i, raw = TRUE)2 -9.028e+00  1.582e+00  -5.707 1.32e-08 ***
## poly(yearfrac, i, raw = TRUE)3  3.923e-03  6.821e-04   5.751 1.03e-08 ***
## poly(yearfrac, i, raw = TRUE)4 -6.817e-07  1.176e-07  -5.795 7.93e-09 ***
## poly(yearfrac, i, raw = TRUE)5          NA          NA          NA          NA
## poly(yearfrac, i, raw = TRUE)6          NA          NA          NA          NA
## poly(yearfrac, i, raw = TRUE)7  2.756e-18  4.649e-19   5.928 3.61e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2504 on 1993 degrees of freedom
## Multiple R-squared:  0.5817, Adjusted R-squared:  0.5806
## F-statistic: 554.3 on 5 and 1993 DF,  p-value: < 2.2e-16
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.24480 -0.15272  0.00679  0.15599  1.01582
##
## Coefficients: (3 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -3.950e+06  7.029e+05  -5.620 2.18e-08 ***
## poly(yearfrac, i, raw = TRUE)1  9.617e+03  1.698e+03   5.663 1.70e-08 ***
## poly(yearfrac, i, raw = TRUE)2 -9.028e+00  1.582e+00  -5.707 1.32e-08 ***

```

```

## poly(yearfrac, i, raw = TRUE)3  3.923e-03  6.821e-04  5.751 1.03e-08 ***
## poly(yearfrac, i, raw = TRUE)4 -6.817e-07  1.176e-07 -5.795 7.93e-09 ***
## poly(yearfrac, i, raw = TRUE)5          NA          NA          NA          NA
## poly(yearfrac, i, raw = TRUE)6          NA          NA          NA          NA
## poly(yearfrac, i, raw = TRUE)7  2.756e-18  4.649e-19  5.928 3.61e-09 ***
## poly(yearfrac, i, raw = TRUE)8          NA          NA          NA          NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2504 on 1993 degrees of freedom
## Multiple R-squared:  0.5817, Adjusted R-squared:  0.5806
## F-statistic: 554.3 on 5 and 1993 DF,  p-value: < 2.2e-16
##
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.24480 -0.15272  0.00679  0.15599  1.01582
##
## Coefficients: (4 not defined because of singularities)
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -3.950e+06  7.029e+05  -5.620 2.18e-08 ***
## poly(yearfrac, i, raw = TRUE)1  9.617e+03  1.698e+03   5.663 1.70e-08 ***
## poly(yearfrac, i, raw = TRUE)2 -9.028e+00  1.582e+00  -5.707 1.32e-08 ***
## poly(yearfrac, i, raw = TRUE)3  3.923e-03  6.821e-04   5.751 1.03e-08 ***
## poly(yearfrac, i, raw = TRUE)4 -6.817e-07  1.176e-07  -5.795 7.93e-09 ***
## poly(yearfrac, i, raw = TRUE)5          NA          NA          NA          NA
## poly(yearfrac, i, raw = TRUE)6          NA          NA          NA          NA
## poly(yearfrac, i, raw = TRUE)7  2.756e-18  4.649e-19   5.928 3.61e-09 ***
## poly(yearfrac, i, raw = TRUE)8          NA          NA          NA          NA
## poly(yearfrac, i, raw = TRUE)9          NA          NA          NA          NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2504 on 1993 degrees of freedom
## Multiple R-squared:  0.5817, Adjusted R-squared:  0.5806
## F-statistic: 554.3 on 5 and 1993 DF,  p-value: < 2.2e-16
##
##
## Call:
## lm(formula = median.temp.anom ~ poly(yearfrac, i, raw = TRUE),
##     data = hadcrut)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.24480 -0.15272  0.00679  0.15599  1.01582
##
## Coefficients: (5 not defined because of singularities)
##
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -3.950e+06  7.029e+05  -5.620 2.18e-08 ***
## poly(yearfrac, i, raw = TRUE)1  9.617e+03  1.698e+03   5.663 1.70e-08 ***

```

```
## poly(yearfrac, i, raw = TRUE)2 -9.028e+00 1.582e+00 -5.707 1.32e-08 ***
## poly(yearfrac, i, raw = TRUE)3 3.923e-03 6.821e-04 5.751 1.03e-08 ***
## poly(yearfrac, i, raw = TRUE)4 -6.817e-07 1.176e-07 -5.795 7.93e-09 ***
## poly(yearfrac, i, raw = TRUE)5 NA NA NA NA
## poly(yearfrac, i, raw = TRUE)6 NA NA NA NA
## poly(yearfrac, i, raw = TRUE)7 2.756e-18 4.649e-19 5.928 3.61e-09 ***
## poly(yearfrac, i, raw = TRUE)8 NA NA NA NA
## poly(yearfrac, i, raw = TRUE)9 NA NA NA NA
## poly(yearfrac, i, raw = TRUE)10 NA NA NA NA
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2504 on 1993 degrees of freedom
## Multiple R-squared: 0.5817, Adjusted R-squared: 0.5806
## F-statistic: 554.3 on 5 and 1993 DF, p-value: < 2.2e-16
```

- 7) If we center or, even better, make the polynomial functions orthogonal to one another, we can avoid the issue in the previous question. Use `poly(x1,?,raw=F)` and step up the polynomial order for time until the p-value for the last coefficient (use `summary()`) is “large”, reporting the single test result for each step in the building process. Then drop back one order and re-fit the model. Report the `effects` plot of the resulting model and describe the estimated trend. Note: aside from access to orthogonal polynomials the `poly` function interfaces with `Anova` and the `effects` package really nicely.

The first coefficient with a “large” p-value was the 11th-order coefficient. Thus, we re-fit the model with a 10th order polynomial. The effects plot of the resulting model is also included.

```
pvals <- numeric(20)
for (i in 1:20){
  lm.poly.1 = lm(median.temp.anom ~ poly(yearfrac,i,raw = FALSE), data = hadcrut)
  pvals[i] = (summary(lm.poly.1)$coefficients[i,4])
}
pvals
```

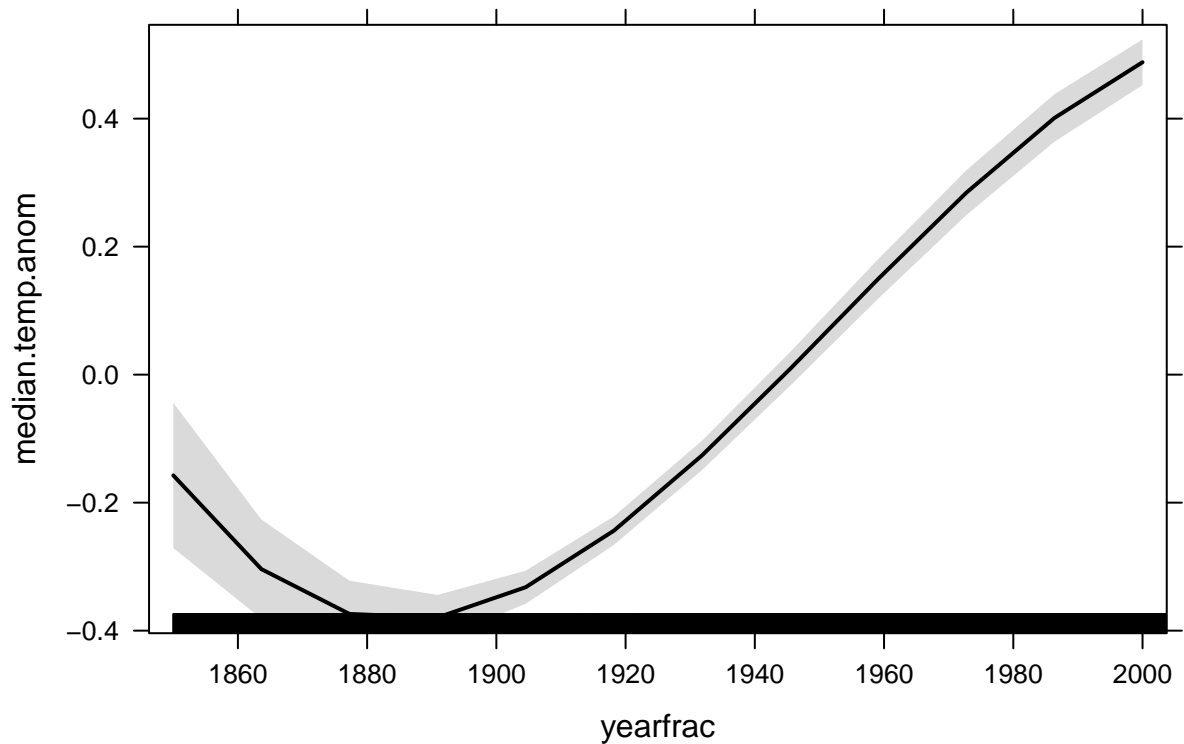
```
## [1] 1.920077e-16 9.531183e-291 5.459973e-98 3.087409e-11 7.386010e-24
## [6] 1.925841e-09 1.294148e-10 6.512222e-22 1.526616e-07 7.239231e-09
## [11] 1.867471e-01 5.514919e-03 5.019406e-01 3.568316e-09 4.328921e-03
## [16] 1.572896e-04 4.041679e-02 2.487755e-01 7.093872e-03 5.932859e-01
```

```
which(pvals > 0.05)
```

```
## [1] 11 13 18 20
```

```
lm.poly.final = lm(median.temp.anom ~ poly(yearfrac,10,raw = FALSE), data = hadcrut)
plot(allEffects(lm.poly.final))
```

yearfrac effect plot

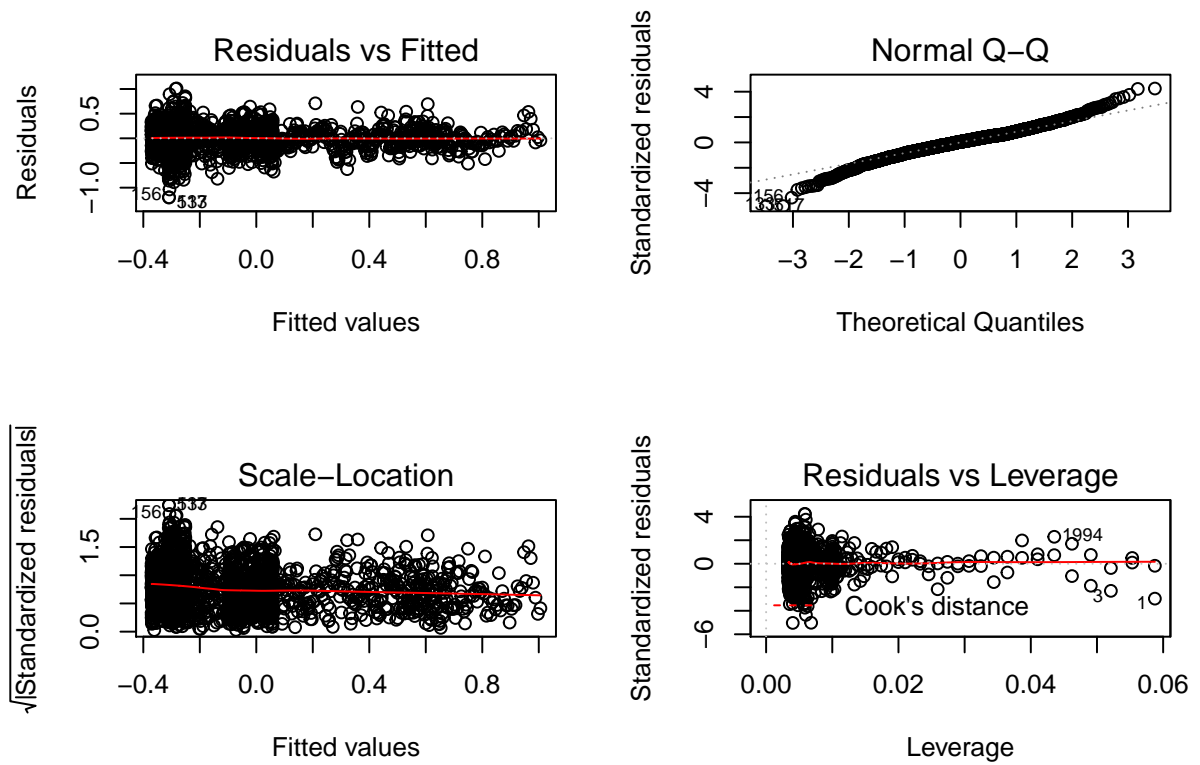


- 8) Check the diagnostic plots from your final model. Does anything improve from the first version. Also plot the residuals vs time and compare that plot to residuals vs fitted.


Based on the plots below, we can see that...

Further, the plot of residuals vs. time shows...

```
par(mfrow = c(2,2))  
plot(lm.poly.final)
```



```
par(mfrow = c(1,1))
plot(hadcrut$yearfrac, resid(lm.poly.final), main = "Residuals vs.Time", xlab = "Time", ylab = "Residuals")
```



9) Run the following code so I can see what version of R you are now using:

Documenting R version

```
sessionInfo()$R.version$nickname
```

```
## [1] "Bug in Your Hair"
```

```
getRversion()
```

```
## [1] '3.3.1'
```

```
sessionInfo()$platform
```

```
## [1] "x86_64-pc-linux-gnu (64-bit)"
```